

Lab 6 : Déploiement K8s d'un système MLOps Churn

Étape 1 : Préparer l'environnement Kubernetes

Instructions :

Minikube installing

```
PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> minikube version
minikube version: v1.37.0
commit: 65318f4cff9c12cc87ec9eb8f4cdd57b25047f3
```

1. Démarrer Minikube (driver Docker) :

```
PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> minikube start --driver=docker --kubernetes-version=v1.28.3
└─ minikube v1.37.0 on Microsoft Windows 11 Pro 10.0.26200.7462 Build 26200.7462
  └─ Using the docker driver based on user configuration
  └─ Using Docker Desktop driver with root privileges
  └─ Starting "minikube" primary control-plane node in "minikube" cluster
  └─ Pulling base image v0.0.48 ...
  └─ Downloading Kubernetes v1.28.3 preload ...
    > gcr.io/k8s-minikube/kicbase...: 561.71 KiB / 488.52 MiB 0.11% 51.72 KiB
```

2. Créer un namespace dédié au lab :

```
PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl create namespace churn-mlops
namespace/churn-mlops created
PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

3. Basculer le contexte courant sur ce namespace :

```
PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl config set-context --current --namespace=churn-mlops
Context "minikube" modified.
PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

4. Vérifier :

```

● PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl get ns
NAME          STATUS   AGE
churn-mllops  Active   0s
default       Active   19m
kube-node-lease Active   19m
kube-public   Active   19m
kube-system   Active   19m
● PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl get pods
No resources found in churn-mllops namespace.
❖ PS C:\Users\anoua\projects\MLOPS\mllops-lab-01>

```

Étape 2 : Préparer l'image Docker de l'API churn

Créer un environnement virtuel avec Python 3.12

Windows / PowerShell et Mettre pip à jour :

```

PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> python --version
Python 3.12.10
PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> py -3.12 -m venv venv_mllops
PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> .\venv_mllops\Scripts\activate
(venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> python -m pip install --upgrade pip
Requirement already satisfied: pip in c:\users\anoua\projects\mllops\mllops-lab-01\venv_mllops\lib\site-packages (25.0.1)
Collecting pip

```

Préciser les dépendances dans `requirements.txt`

(afin d'éviter toute incompatibilité entre entraînement et prédiction)

Créer ou éditer le fichier

et y mettre :

```

fastapi      # Framework web rapide pour créer des APIs Python
uvicorn[standard] # Serveur ASGI pour exécuter FastAPI
pydantic     # Validation et typage des données
scikit-learn==1.7.2 # Entraînement et inférence des modèles ML
pandas==2.2.3    # Manipulation de données tabulaires

```

```
numpy==2.1.3      # Calcul numérique  
joblib==1.4.2     # Sérialisation et chargement des modèles
```

Installer les dépendances :

```
(venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> pip install -r requirements.txt  
collecting fastapi (from -r requirements.txt (line 1))  
  Using cached fastapi-0.128.0-py3-none-any.whl.metadata (30 kB)  
collecting pydantic (from -r requirements.txt (line 3))  
  Using cached pydantic_2.12.5_py3-none-any.whl.metadata (29 kB)
```

Cette étape est obligatoire pour :

éviter les erreurs de compatibilité `scikit-learn`, garantir que le modèle chargé par l'API correspond exactement aux versions installées.

- éviter les erreurs de compatibilité `scikit-learn`,
- garantir que le modèle chargé par l'API correspond exactement aux versions installées.

Étape 3 : Créer le dossier des manifests Kubernetes

Instructions :Sous PowerShell :

```
New-Item -ItemType Directory -Name k8s
```

```
(venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> New-Item -ItemType Directory -Name k8s  
  
Directory: C:\Users\anoua\projects\MLOPS\mllops-lab-01  
  
Mode                LastWriteTime          Length Name  
----                -----          -----   -----  -----  
d-----  1/15/2026  3:21 PM           0    k8s
```

Vérifier :

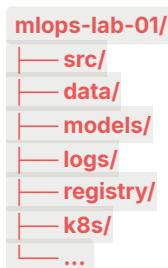
```
Get-ChildItem
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> Get-ChildItem
```

```
Directory: C:\Users\anoua\projects\MLOPS\mlops-lab-01
```

Mode	LastWriteTime	Length	Name
d----	1/4/2026 12:05 PM		.dvc
d----	1/4/2026 3:06 PM		.github
d----	1/4/2026 1:06 PM		data
d----	1/4/2026 12:19 PM		dvc_storage
d----	1/15/2026 3:21 PM		k8s
d----	1/3/2026 10:18 PM		logs
d----	1/5/2026 2:53 AM		models

Tu dois voir :



Étape 4 : Construire l'image Docker (tag versionné)

Instructions : Se placer dans le dossier du projet :

```
cd ./mlops-lab-01
```

Construire l'image Docker avec un tag versionné (jamais `latest`) :

```
# Construire l'image du projet avec un tag de version (v1, v2, ...)
docker build -t churn-api:v1 .
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> docker build -t churn-api:v1 .
[+] Building 3.2s (1/2)
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 501B
=> [internal] load metadata for docker.io/library/python:3.11-slim
               docker:desktop-linux
               0.1s
               0.0s
               3.1s
```

Vérifier l'image localement :

```
# PowerShell
docker images | Select-String churn-api
```

```
(venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> docker images | Select-String churn-api
WARNING: This output is designed for human readability. For machine-readable output, please use --format.

churn-api:latest          15c3c801b7f6      516MB        0B
churn-api:v1               78dd3690aa4a     852MB        0B
```

Étape 5 : Charger explicitement l'image dans Minikube

Instructions : Sauvegarder l'image Docker :

```
docker save churn-api:v1 -o churn-api_v1.tar
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> docker save churn-api:v1 -o churn-api_v1.tar
✖ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

Charger l'image dans Minikube :

```
minikube image load churn-api_v1.tar
```

```
(venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> minikube image load churn-api_v1.tar
(venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

Vérifier que l'image est disponible dans Minikube :

```
# PowerShell
minikube image ls | Select-String churn-api
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> minikube image ls | select-string churn-api
  docker.io/library/churn-api:v1

○ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

Étape 6 : Deployment Kubernetes pour l'API churn

Instructions :

1. Créer le fichier :

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> New-Item k8s/deployment.yaml

Directory: C:\Users\anoua\projects\MLOPS\mlops-lab-01\k8s

Mode                LastWriteTime     Length Name
----                -----          ---- 
-a--- 1/15/2026 3:44 PM           0 deployment.yaml
```

2. Coller le contenu suivant dans `k8s/deployment.yaml` :

```
① README.md X  Ⓜ requirements.txt M  ! deployment.yaml U X

k8s > ! deployment.yaml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: churn-api
5  spec:
6    replicas: 2
7    selector:
8      matchLabels:
9        app: churn-api
10   template:
11     metadata:
12       labels:
13         app: churn-api
14   spec:
15     containers:
16       - name: api
17         image: churn-api:v1 # v1 presente le tag de l'image cible
18         ports:
19           - containerPort: 8000
```

3. Appliquer le manifest :

```
kubectl apply -f k8s/deployment.yaml
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl apply -f k8s/deployment.yaml
deployment.apps/churn-api created
✿ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

Suivre le rollout :

kubectl rollout status deployment churn-api

```
• (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl rollout status deployment churn-api  
deployment "churn-api" successfully rolled out  
❖ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> █
```

Vérifier :

```
kubectl get pods -l app=churn-api -o wide
```

```
• (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl get pods -l app=churn-api -o wide
NAME          READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
churn-api-7d686dfbf7-lc7d4   1/1    Running   0          53s    10.244.0.4   minikube   <none>   <none>
churn-api-7d686dfbf7-z278p   1/1    Running   0          53s    10.244.0.5   minikube   <none>   <none>
❖ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

Étape 7 : Exposer l'API via un Service NodePort

Instructions :Créer le fichier :

New-Item k8s/service.yaml

```
(venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> New-Item k8s/service.yaml  
  
Directory: C:\Users\anoua\projects\MLOPS\mlops-lab-01\k8s  
  
Mode          LastWriteTime        Length Name  
----          -----          ---- -  
-a---  1/15/2026 3:57 PM           0 service.yaml
```

Coller le contenu suivant dans k8s/service.yaml :

```
k8s > ! service.yaml
 1 apiVersion: v1
 2 kind: Service
 3 metadata:
 4   name: churn-api-service
 5 spec:
 6   type: NodePort
 7   selector:
 8     app: churn-api
 9   ports:
10     - port: 80
11       targetPort: 8000
12       nodePort: 30080
```

Appliquer :

```
kubectl apply -f k8s/service.yaml
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl apply -f k8s/service.yaml
>>
service/churn-api-service created
* (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

Visualiser la création du Service :

```
kubectl get svc churn-api-service
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl get svc churn-api-service
  NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
churn-api-service   NodePort    10.96.57.157    <none>        80:30080/TCP   19s
* (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

Afficher les détails du Service :

```
kubectl describe svc churn-api-service
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl describe svc churn-api-service
Name:           churn-api-service
Namespace:      default
Labels:         <none>
Annotations:   <none>
Selector:       app=churn-api
Type:          NodePort
IP Family Policy: SingleStack
IP Families:   IPv4
IP:            10.96.57.157
IPs:           10.96.57.157
Port:          <unset>  80/TCP
TargetPort:    8000/TCP
NodePort:      <unset>  30080/TCP
Endpoints:     10.244.0.5:8000,10.244.0.4:8000
Session Affinity: None
External Traffic Policy: Cluster
Internal Traffic Policy: Cluster
Events:        <none>

```

Ouvrir l'accès à la communication avec le cluster (Windows / Minikube driver Docker) :

`kubectl port-forward svc/churn-api-service 30080:80`

```
Forwarding from [::]:30080 -> 8000
Forwarding from 127.0.0.1:30080 -> 8000
Forwarding from [::1]:30080 -> 8000
```

Tester l'API, via Postman :

The screenshot shows the Postman application interface. A POST request is being made to the URL `http://127.0.0.1:30080/predict`. The "Body" tab is active, and the raw JSON payload is displayed in the text area:

```

1 {
2   "tenure_months": 48,
3   "num_complaints": 0,
4   "avg_session_minutes": 60,
5   "plan_type": "premium",
6   "region": "EU",
7   "request_id": "req-safe"
8 }

```

résultat :

```
1 {
2     "request_id": "req-safe",
3     "model_version": "churn_model_v1_20260115_151424.joblib",
4     "prediction": 0,
5     "probability": 0.139973,
6     "latency_ms": 10.362,
7     "features": {
8         "tenure_months": 48,
9         "num_complaints": 0,
10        "avg_session_minutes": 60.0,
11        "plan_type": "premium",
12        "region": "EU"
13    },
14    "ts": 1768491259
15 }
```

Étape 8 : Injecter la configuration MLOps via ConfigMap

Instructions :Créer :

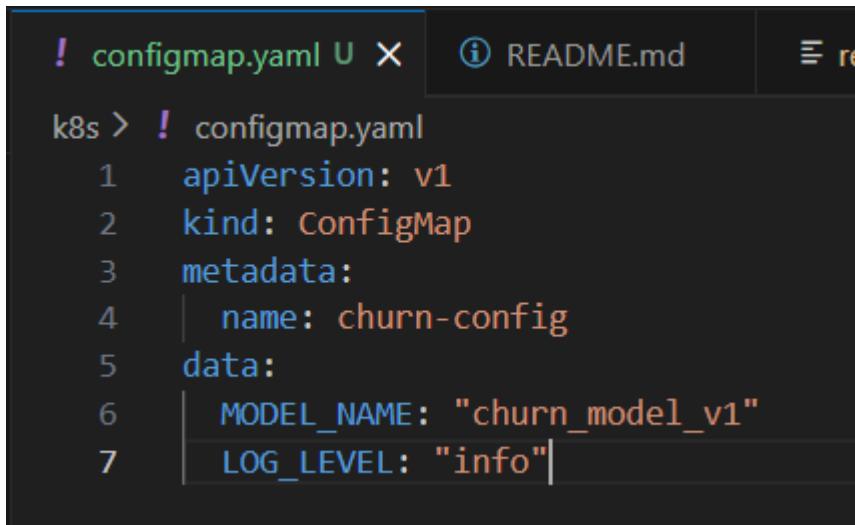
New-Item k8s/configmap.yaml

● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> New-Item k8s/configmap.yaml

Directory: C:\Users\anoua\projects\MLOPS\mlops-lab-01\k8s

Mode	LastWriteTime	Length	Name
----	-----	-----	0 configmap.yaml

Coller :



```
! configmap.yaml U X ⓘ README.md ⌂ re
k8s > ! configmap.yaml
1   apiVersion: v1
2   kind: ConfigMap
3   metadata:
4     name: churn-config
5   data:
6     MODEL_NAME: "churn_model_v1"
7     LOG_LEVEL: "info"
```

Appliquer :

```
kubectl apply -f k8s/configmap.yaml
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl apply -f k8s/configmap.yaml
configmap/churn-config created
❖ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

Visualiser la création du ConfigMap :

```
kubectl get configmap churn-config
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl get configmap churn-config
NAME      DATA   AGE
churn-config   2    21s
```

Afficher le contenu du ConfigMap :

```
kubectl describe configmap churn-config
```

```
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl describe configmap churn-config
Name:          churn-config
Namespace:     default
Namespace:     default
Labels:        <none>
Namespace:     default
Labels:        <none>
Annotations:   <none>

Data
====
LOG_LEVEL:
-----
info

MODEL_NAME:
-----
churn_model_v1

BinaryData
-----
Events:  <none>
(venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01>
```

Modifier `k8s/deployment.yaml`

pour injecter ces variables dans le conteneur : Dans la section `containers: -`
`name: api`, ajouter :

```

! configmap.yaml U | ① README.md | ⌂ requirements.txt M | ! deployment.yaml U
k8s > ! deployment.yaml
  5   spec:
    7     selector:
  10   template:
  11     metadata:
  12       labels:
  13         app: churn-api
  14     spec:
  15       containers:
  16         - name: api
  17           image: churn-api:v1 # v1 presente le tag de l'image cible
  18           ports:
  19             - containerPort: 8000
  20           env:
  21             - name: MODEL_NAME
  22               valueFrom:
  23                 configMapKeyRef:
  24                   name: churn-config
  25                   key: MODEL_NAME
  26             - name: LOG_LEVEL
  27               valueFrom:
  28                 configMapKeyRef:
  29                   name: churn-config
  30                   key: LOG_LEVEL

```

Réappliquer le Deployment :

```
kubectl apply -f k8s/deployment.yaml
```

```

EVENTS:  <none>
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl apply -f k8s/deployment.yaml
deployment.apps/churn-api configured
❖ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

Suivre le rollout :

```

kubectl rollout restart deployment churn-api
kubectl rollout status deployment churn-api

```

```
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl rollout restart deployment churn-api
deployment.apps/churn-api restarted
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl rollout status deployment churn-api
deployment "churn-api" successfully rolled out
❖ (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> s
```

Vérifier que les variables sont bien injectées dans un Pod :

```
kubectl exec -it deploy/churn-api -- printenv MODEL_NAME
kubectl exec -it deploy/churn-api -- printenv LOG_LEVEL
```

```
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl exec -it deploy/churn-api -- printenv MODEL_NAME
churn_model_v1
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl exec -it deploy/churn-api -- printenv LOG_LEVEL
info
```

Étape 9 : Gérer les secrets (MONITORING_TOKEN)

Instructions :Encoder une valeur simple en base64 (exemple "abc123") :
Depuis PowerShell :

```
[Convert]::ToString([Text.Encoding]::UTF8.GetBytes("abc123"))
```

```
(venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> [Convert]::ToString([Text.Encoding]::UTF8.GetBytes("abc123"))
YWJjMTIz
(venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01>
```

Garder la valeur obtenue, par ex : **YWJjMTIz** Créer le fichier :

```
New-Item k8s/secret.yaml
```

```
(venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> New-Item k8s/secret.yaml

Directory: C:\Users\anoua\projects\MLOPS\mllops-lab-01\k8s

Mode                LastWriteTime          Length Name
----                -----          ---- 
-a--- 1/15/2026 5:04 PM           0 secret.yaml
```

Coller :

```
k8s > ! secret.yaml
 1  apiVersion: v1
 2  kind: Secret
 3  metadata:
 4    name: churn-secret
 5  type: Opaque
 6  data:
 7    MONITORING_TOKEN: "YWJjMTIz"
```

Appliquer et Visualiser la création du Secret :

```
kubectl apply -f k8s/secret.yaml
kubectl get secret churn-secret
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl apply -f k8s/secret.yaml
secret/churn-secret created
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl get secret churn-secret
  NAME      TYPE      DATA   AGE
churn-secret  Opaque     1    27s
```

Afficher les détails (sans afficher les valeurs en clair) :

```
kubectl describe secret churn-secret
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl describe secret churn-secret
Name:          churn-secret
Namespace:     default
Labels:        <none>
Annotations:  <none>

Type:  Opaque

Data
====
MONITORING_TOKEN:  6 bytes
◆ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

Ajouter la variable d'environnement dans `k8s/deployment.yaml` (dans `env:` déjà présent) :

```

k8s > ! deployment.yaml
26      - name: LOG_LEVEL
27          valueFrom:
28              configMapKeyRef:
29                  name: churn-config
30                  key: LOG_LEVEL
31      - name: MONITORING_TOKEN
32          valueFrom:
33              secretKeyRef:
34                  name: churn-secret
35                  key: MONITORING_TOKEN

```

Réappliquer ET Vérifier le redémarrage et l'état des Pods :

```
kubectl apply -f k8s/deployment.yaml
```

```
kubectl get pods
```

```

● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl apply -f k8s/deployment.yaml
deployment.apps/churn-api configured
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl get pods
  NAME           READY   STATUS    RESTARTS   AGE
  churn-api-54b6c6d49b-d6scb   0/1     Terminating   0          40m
  churn-api-54b6c6d49b-jk76f   1/1     Terminating   0          40m
  churn-api-5cc7f85545-gw2rd   1/1     Running    0          4s
  churn-api-5cc7f85545-qlbpq   1/1     Running    0          6s
◆ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>

```

Vérifier que la variable d'environnement est bien injectée dans un Pod :

```
kubectl exec -it deploy/churn-api -- printenv MONITORING_TOKEN
```

```

(venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl exec -it deploy/churn-api -- printenv M
● ONITORING_TOKEN
abc123
◆ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>

```

Étape 10 : Mise en place des endpoints de santé et des probes Kubernetes pour l'API Churn

Instructions : Ajouter les endpoints nécessaires aux probes Kubernetes dans l'API.

Ouvrir le fichier contenant la définition de l'API FastAPI (par exemple

[src/api.py](#)).

Ajouter les endpoints suivants :

```
@app.get("/health")
def health() → dict[str, Any]:
    """
```

Endpoint de santé de l'API.

Vérifie simplement qu'un modèle courant est bien configuré.

Retour

dict

- status : "ok" ou "error"
- current_model : nom du modèle courant (si OK)
- detail : message d'erreur (si error)

"""

try:

```
    model_name = get_current_model_name()
    return {"status": "ok", "current_model": model_name}
```

```
except Exception as exc: # pragma: no cover - simple endpoint de debu
g
```

```
    return {"status": "error", "detail": str(exc)}
```

```
@app.get("/startup")
```

```
def startup() → dict[str, Any]:
```

"""

Endpoint utilisé par Kubernetes startupProbe.

L'application est considérée comme démarrée UNIQUEMENT si :

- le registry existe,
- le fichier current_model.txt existe,
- le fichier n'est pas vide.

"""

```
if not REGISTRY_DIR.exists():
```

```
    raise HTTPException(
```

```

        status_code=503,
        detail="Registry non monté (PVC absent ou incorrect).",
    )

if not CURRENT_MODEL_PATH.exists():
    raise HTTPException(
        status_code=503,
        detail="Aucun modèle courant. Lancer train.py (avec gate) d'abor
d.",
    )

name = CURRENT_MODEL_PATH.read_text(encoding="utf-8").strip()
if not name:
    raise HTTPException(
        status_code=503,
        detail="current_model.txt vide.",
    )

return {
    "status": "ok",
    "current_model": name,
}

@app.get("/ready")
def ready() → dict[str, Any]:
    try:
        model_name = get_current_model_name()
        return {"status": "ready", "current_model": model_name}
    except Exception as exc:
        raise HTTPException(status_code=503, detail=str(exc))

```

```

src > ⚡ api.py > 📄 health
217
218     @app.get("/health")
219     def health() -> dict[str, Any]:
220         """
221             Endpoint de santé de l'API.
222
223
224             Vérifie simplement qu'un modèle courant est bien configuré.
225
226
227             Retour
228             -----
229             dict
230                 - status : "ok" ou "error"
231                 - current_model : nom du modèle courant (si OK)
232                 - detail : message d'erreur (si error)
233             """
234     try:
235         model_name = get_current_model_name()
236         return {"status": "ok", "current_model": model_name}
237     except Exception as exc: # pragma: no cover - simple endpoint de debug
238         return {"status": "error", "detail": str(exc)}
239
240
241
242     @app.get("/startup")
243     def startup() -> dict[str, Any]:
244         """
245             Endpoint utilisé par Kubernetes startupProbe.
246
247             L'application est considérée comme démarrée UNIQUEMENT si :
248                 - le registry existe,
249                 - le fichier current_model.txt existe

```

Sauvegarder le fichier.

Reconstruire l'image Docker de l'API :

```
docker build-t churn-api:v1 .
```

```

❖ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> docker build -t churn-api:v1 .
[+] Building 46.1s (8/9)                                            docker:desktop-linux
=> [internal] load build definition from Dockerfile                  0.1s
=> => transferring dockerfile: 501B                                  0.0s
=> [internal] load metadata for docker.io/library/python:3.12-slim   1.5s
=> [internal] load .dockerignore                                    0.0s
=> => transferring context: 2B                                     0.0s
=> [1/5] FROM docker.io/library/python:3.12-slim@sha256:d75c4b6cdd039ae966a34cd3ccab9e0e5f7299280ad  0.0s
=> [internal] load build context                                    41.2s

```

Exporter l'image Docker :

```
docker save churn-api:v1 -o churn-api_v1.tar
```

Charger l'image dans Minikube :

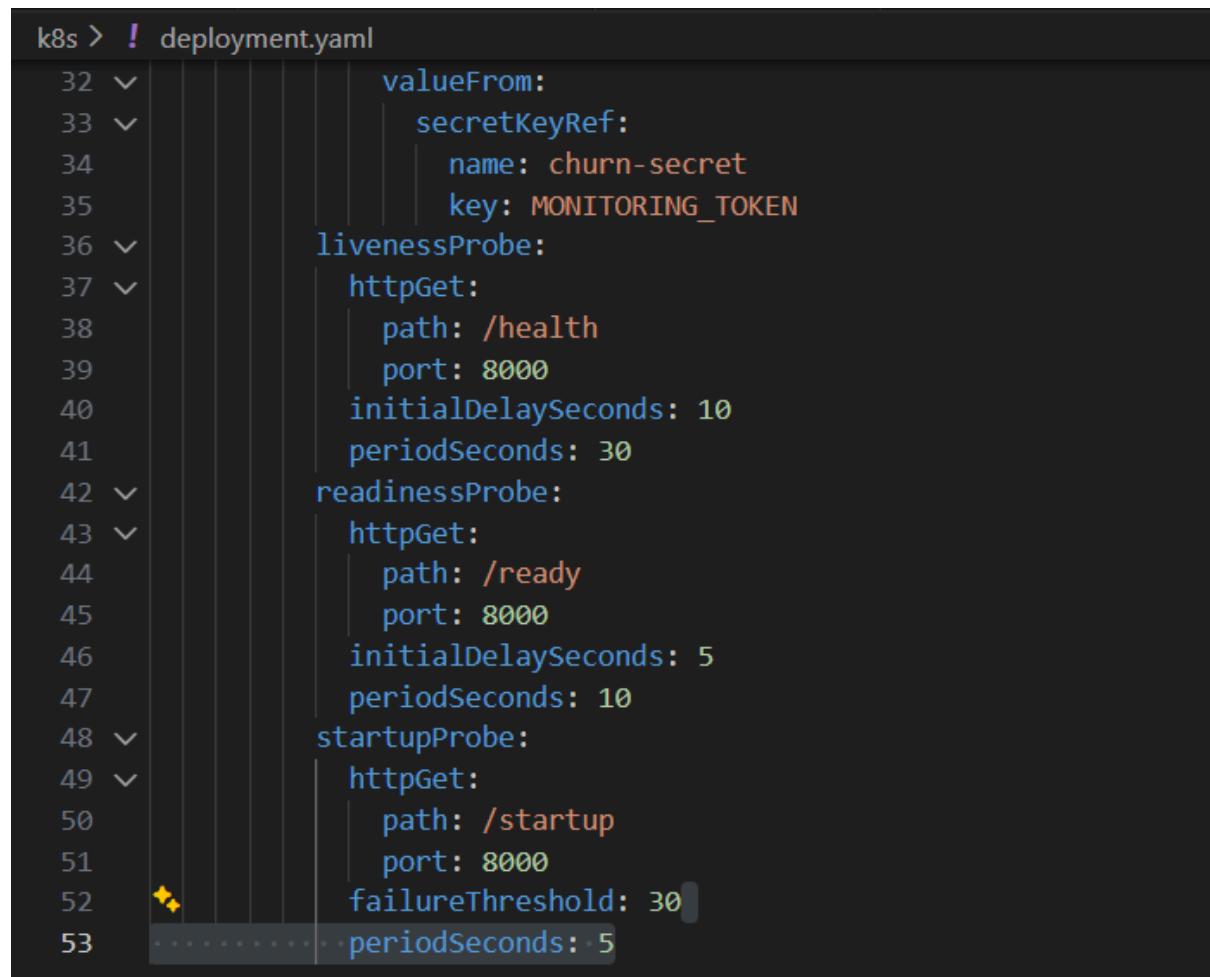
```
minikube image load churn-api_v1.tar
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> docker save churn-api:v1 -o churn-api_v1.tar
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> minikube image load churn-api_v1.tar
○ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> [REDACTED]
```

Étape 11 : Ajouter les probes (liveness / readiness / startup)

Instructions :

1. Dans `k8s/deployment.yaml`, compléter le conteneur :



```
k8s > ! deployment.yaml
32     valueFrom:
33         secretKeyRef:
34             name: churn-secret
35             key: MONITORING_TOKEN
36     livenessProbe:
37         httpGet:
38             path: /health
39             port: 8000
40             initialDelaySeconds: 10
41             periodSeconds: 30
42     readinessProbe:
43         httpGet:
44             path: /ready
45             port: 8000
46             initialDelaySeconds: 5
47             periodSeconds: 10
48     startupProbe:
49         httpGet:
50             path: /startup
51             port: 8000
52             failureThreshold: 30
53             periodSeconds: 5
```

2. Réappliquer :

```
kubectl apply -f k8s/deployment.yaml  
kubectl describe pod -l app=churn-api
```

● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> **kubectl** apply -f k8s/deployment.yaml
deployment.apps/churn-api configured
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> **kubectl** describe pod -l app=churn-api
Name: churn-api-5cc7f85545-gw2rd
Namespace: default
Priority: 0
Service Account: default
Node: minikube/192.168.49.2
Start Time: Thu, 15 Jan 2026 17:35:14 +0100
Labels: app=churn-api
pod-template-hash=5cc7f85545
Annotations: kubectl.kubernetes.io/restartedAt: 2026-01-15T16:55:58+01:00
Status: Running
IP: 10.244.0.13
IPs:
IP: 10.244.0.13
Controlled By: ReplicaSet/churn-api-5cc7f85545
Containers:
api:
 Container ID: docker://643e981d5ba14eca016a319b3f831c609cf54ddcf9ddd18924e2309b2a157e4
 Image: churn-api:v1

Redéployer l'application dans le cluster Kubernetes :

```
kubectl rollout restart deployment churn-api  
kubectl rollout status deployment churn-api
```

● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> **kubectl** rollout restart deployment churn-api
deployment.apps/churn-api restarted
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> **kubectl** rollout status deployment churn-api
Waiting for deployment "churn-api" rollout to finish: 1 out of 2 new replicas have been updated...
Waiting for deployment "churn-api" rollout to finish: 1 out of 2 new replicas have been updated...
Waiting for deployment "churn-api" rollout to finish: 1 out of 2 new replicas have been updated...
Waiting for deployment "churn-api" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "churn-api" rollout to finish: 1 old replicas are pending termination...
deployment "churn-api" successfully rolled out
❖ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>

Vérifier l'état des Pods :

```
kubectl get pods
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl get pods
  NAME           READY   STATUS    RESTARTS   AGE
  churn-api-f6fd88cc4-7ll5j   1/1     Running   0          3m23s
  churn-api-f6fd88cc4-tjqdf   1/1     Running   0          2m19s
◆ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

Objectif :

Kubernetes redémarre un Pod si `/health` ne répond plus, et n'envoie du trafic qu'aux Pods ready.

Étape 12 : Volume persistant pour registry + logs

Instructions :Créer le PVC (PersistentVolumeClaim) : ce volume persistant permet de conserver les fichiers du modèle et des logs même si les Pods sont recréés.

Créer le fichier :

New-Item k8s/pvc.yaml

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> New-Item k8s/pvc.yaml

Directory: C:\Users\anoua\projects\MLOPS\mlops-lab-01\k8s

Mode                LastWriteTime         Length Name
----                -----          0 pvc.yaml
```

Coller :

```

k8s > ! pvc.yaml
 1  apiVersion: v1
 2  kind: PersistentVolumeClaim
 3  metadata:
 4    name: churn-storage
 5  spec:
 6    accessModes:
 7      - ReadWriteOnce
 8    resources:
 9      requests:
10        storage: 5Gi

```

Appliquer :

```
kubectl apply -f k8s/pvc.yaml
```

```

● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl apply -f k8s/pvc.yaml
persistentvolumeclaim/churn-storage created
◆ (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01>

```

Vérifier que le PVC est bien créé et associé à un volume :

```
kubectl get pvc
```

```

● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl get pvc
  NAME      STATUS   VOLUME                                     CAPACITY  ACCESS MODES  STORAGECLASS
  AGE
  churn-storage   Bound   pvc-bdf3ed8d-b164-4ee7-b2cd-4725bccfd7c   5Gi       RWO          standard
  76s
◆ (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01>

```

Créer et exécuter le Job d'entraînement : ce Job initialise le contenu du PVC en entraînant un premier modèle et en écrivant les artefacts dans [/app/registry](#) .

Créer le fichier :

```
New-Item k8s/job-train.yaml
```

```

● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> New-Item k8s/job-train.yaml

Directory: C:\Users\anoua\projects\MLOPS\mlops-lab-01\k8s

Mode                LastWriteTime          Length Name
----                -----          0 job-train.yaml

-a---- 1/15/2026 7:39 PM

```

Coller :

```

k8s > ! job-train.yaml
 1  apiVersion: batch/v1
 2  kind: Job
 3  metadata:
 4    name: churn-train
 5  spec:
 6    backoffLimit: 1
 7    template:
 8      spec:
 9        restartPolicy: Never
10        volumes:
11          - name: churn-volume
12            persistentVolumeClaim:
13              claimName: churn-storage
14        containers:
15          - name: train
16            image: churn-api:v1
17            command: ["python", "src/train.py"]
18            volumeMounts:
19              - name: churn-volume
20                mountPath: /app/models
21                subPath: models
22              - name: churn-volume
23                mountPath: /app/registry
24                subPath: registry

```

Appliquer le Job :

`kubectl apply -f k8s/job-train.yaml`

```
(venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl apply -f k8s/job-train.yaml
job.batch/churn-train created
(venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01>
```

Attendre la fin du Job :

```
kubectl wait --for=condition=complete job/churn-train
```

```
(venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl wait --for=condition=complete job/churn-train
● job.batch/churn-train condition met
❖ (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01>
```

Monter le PVC dans le Deployment : l'API doit lire le modèle courant et écrire ses logs dans le même stockage persistant.

Modifier `k8s/deployment.yaml` pour ajouter le volume dans `spec.template.spec`

Dans la section `containers: - name: api`, monter ce volume à deux emplacements distincts :

- `/app/registry` pour stocker et lire les modèles,
- `/app/logs` pour écrire les logs applicatifs

```

k8s > ! deployment.yaml
 45   |     port: 8000
 46   |     initialDelaySeconds: 5
 47   |     periodSeconds: 10
 48   |     startupProbe:
 49   |       httpGet:
 50   |         path: /startup
 51   |         port: 8000
 52   |     failureThreshold: 30
 53   |     periodSeconds: 5
 54   |     volumeMounts:
 55   |       - name: churn-volume
 56   |         mountPath: /app/registry
 57   |         subPath: registry
 58   |       - name: churn-volume
 59   |         mountPath: /app/models
 60   |         subPath: models
 61   |       - name: churn-volume
 62   |         mountPath: /app/logs
 63   |         subPath: logs
 64   |     volumes:
 65   |       - name: churn-volume
 66   |         persistentVolumeClaim:
 67   |           claimName: churn-storage

```

Réappliquer :

`kubectl apply -f k8s/deployment.yaml`

```

● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl apply -f k8s/deployment.yaml
deployment.apps/churn-api configured
❖ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>

```

Vérifier que les Pods redémarrent correctement :

`kubectl get pods`

```

● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl get pods
  NAME          READY   STATUS    RESTARTS   AGE
churn-api-7655fd649b-tg6fz   1/1     Running   0          5m50s
churn-api-7655fd649b-vzd4l   1/1     Running   0          6m50s
churn-train-6kzf6            0/1     Completed  0          9m30s
❖ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>

```

Vérifier que les dossiers sont bien accessibles depuis un Pod :

```
kubectl exec -it deploy/churn-api -- ls /app/registry  
kubectl exec -it deploy/churn-api -- ls /app/logs
```

```
(venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl exec -it deploy/churn-api -- ls /app/re  
gistry  
current_model.txt metadata.json  
(venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl exec -it deploy/churn-api -- ls /app/lo  
gs  
(venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>
```

Étape 13 : NetworkPolicy

Instructions :

1. Créer le fichier :

```
New-Item k8s/networkpolicy.yaml
```

```
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> New-Item k8s/networkpolicy.yaml  
  
Directory: C:\Users\anoua\projects\MLOPS\mlops-lab-01\k8s  
  
Mode                LastWriteTime          Length Name  
----                -----          ---- -  
-a----       1/15/2026    7:54 PM           0 networkpolicy.yaml
```

2. Coller le contenu suivant :

```

k8s > ! networkpolicy.yaml
  1  apiVersion: networking.k8s.io/v1
  2  kind: NetworkPolicy
  3  metadata:
  4    name: allow-internal-services
  5  spec:
  6    podSelector:
  7      matchLabels:
  8        app: churn-api
  9    policyTypes:
 10      - Ingress
 11    ingress:
 12      - from:
 13        - podSelector: {}
 14      ports:
 15        - port: 8000
 16          protocol: TCP

```

3. Appliquer la configuration :

```
kubectl apply -f k8s/networkpolicy.yaml
kubectl get networkpolicy
```

```
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl apply -f k8s/networkpolicy.yaml
networkpolicy.networking.k8s.io/allow-internal-services created
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl get networkpolicy
NAME                POD-SELECTOR   AGE
allow-internal-services   app=churn-api   4s
❖ (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01>
```

Étape 14 : Vérifications finales

Instructions :

1. Vérifier les Pods et les Services :

```
kubectl get pods -l app=churn-api
kubectl get svc
```

```

● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl get pods -l app=churn-api
  NAME                READY   STATUS    RESTARTS   AGE
  churn-api-7655fd649b-tg6fz   1/1     Running   0          12m
  churn-api-7655fd649b-vzd4l   1/1     Running   0          13m
● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl get svc
  NAME            TYPE        CLUSTER-IP      EXTERNAL-IP    PORT(S)        AGE
  churn-api-service   NodePort    10.100.129.224 <none>       80:30080/TCP   3h22m
  kubernetes       ClusterIP   10.96.0.1     <none>       443/TCP       29h
◆ (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01>

```

2. Tester l'endpoint /health :

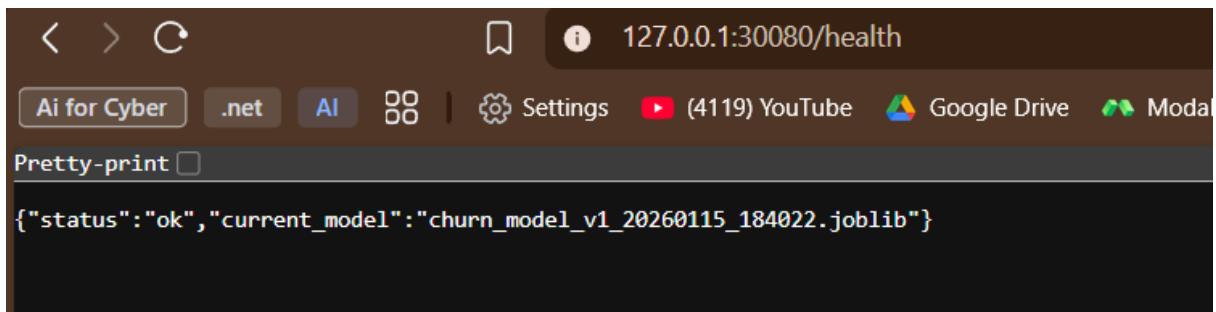
kubectl port-forward svc/churn-api-service 30080:80

```

● (venv_mlops) PS C:\Users\anoua\projects\MLOPS\mlops-lab-01> kubectl port-forward svc/churn-api-service 30080:80
  Forwarding from 127.0.0.1:30080 -> 8000
  Forwarding from [::1]:30080 -> 8000

```

GET http://127.0.0.1:30080/health



3. Envoyer quelques requêtes /predict.

POST http://127.0.0.1:30080/predict

Params Authorization Headers (9) Body **JSON** Pre-request Script Tests Settings

Body (raw JSON)

```
1 {  
2     "tenure_months": 55,  
3     "num_complaints": 1,  
4     "avg_session_minutes": 36,  
5     "plan_type": "premium",  
6     "region": "EU",  
7     "request_id": "req-safe"  
8 }
```

```
1 {  
2     "request_id": "req-safe",  
3     "model_version": "churn_model_v1_20260115_184022.joblib",  
4     "prediction": 0,  
5     "probability": 0.22183,  
6     "latency_ms": 11.417,  
7     "features": {  
8         "tenure_months": 55,  
9         "num_complaints": 1,  
0         "avg_session_minutes": 36.0,  
1         "plan_type": "premium",  
2         "region": "EU"  
3     },  
4     "ts": 1768503534  
5 }
```

4. Lister les Pods pour choisir un Pod churn :

```
kubectl get pods -l app=churn-api
```

```
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl get pods -l app=churn-api  
NAME          READY   STATUS    RESTARTS   AGE  
churn-api-7655fd649b-tg6fz  1/1     Running   0          16m  
churn-api-7655fd649b-vzd4l  1/1     Running   0          17m  
◆ (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01>
```

5. Exécuter la détection de drift dans le Pod :

```
kubectl exec -it <nom_du_pod> -- python src/monitor_drift.py
```

problem :

```
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl exec -it churn-api-7655fd649b-tg6fz -c api -- python src/monitor_drift.py
Traceback (most recent call last):
  File "/app/src/monitor_drift.py", line 193, in <module>
    main()
  File "/app/src/monitor_drift.py", line 81, in main
    raise FileNotFoundError
FileNotFoundError: train_stats.json introuvable. Lancer d'abord prepare_data.py pour générer les statistiques.
command terminated with exit code 1
◆ (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> █
```

- Result: Drift check completed with 1 recent request; no drift detected.
- Files present: train_stats.json, processed.csv, logs used from predictions.log.
- PVC mounts: Verified in the pod under /app/registry.

```
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> $pod = "churn-api-7655fd649b-tg6fz"
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl exec -it $pod -c api -- python src/prepare_data.py
[OK] Fichier prétraité généré : /app/data/processed.csv
[OK] Statistiques d'entraînement générées : /app/registry/train_stats.json
HElloo
cheek
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl exec -it $pod -c api -- ls -la /app/registry
total 24
drwxrwxrwx 2 root root 4096 Jan 15 19:07 .
drwxr-xr-x 1 root root 4096 Jan 15 16:39 ..
-rw-r--r-- 1 root root 37 Jan 15 18:40 current_model.txt
-rw-r--r-- 1 root root 409 Jan 15 18:40 metadata.json
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl exec -it $pod -c api -- python src/monitor_drift.py
● (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> kubectl exec -it $pod -c api -- python src/monitor_drift.py
    === Drift check sur 1 requêtes récentes ===
    - tenure_months: mean_prod=55.000 | mean_train=30.246 | z=1.458
    - num_complaints: mean_prod=1.000 | mean_train=1.174 | z=0.157
    - avg_session_minutes: mean_prod=36.000 | mean_train=35.124 | z=0.074
    Résultat : aucun drift détecté.
◆ (venv_mllops) PS C:\Users\anoua\projects\MLOPS\mllops-lab-01> █
```